

Development and evaluation of a Kubernetes cluster simulator based on Batsim

Presented by: Théo Larue

Supervised by: Olivier Richard & Michael Mercier

Université Grenoble Alpes

August 31, 2020



Table of contents

- 1 Introduction
- 2 Literature review
- 3 Integrating Kubernetes schedulers to Batsim
- 4 Study of the simulator
- 5 Discussion and future work


Introduction

Resource and Jobs Management System

The RJMS is at the core of the cluster.

Examples of RJMS


- OAR
- SLURM
- HadoopYARN
- Apache Mesos




Infrastructure



Resource and Job Management System



Operating System



Network



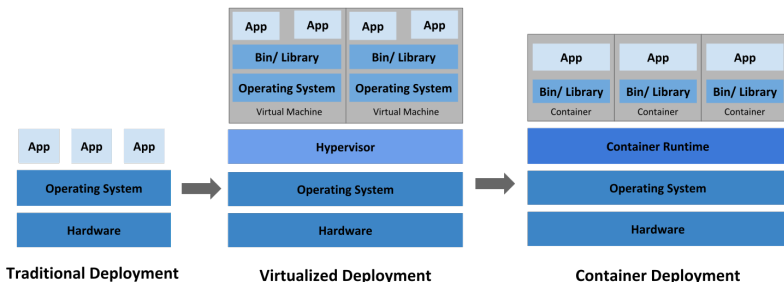
Hardware



Kubernetes


Kubernetes in a nutshell

- Open source resource manager for containerized applications
- About 2M lines of code
- 2.8k contributors



source: <https://kubernetes.io/docs/>

A component of the RJMS: the scheduler



Scheduling is the act of allocating tasks to resources.

Numerous factors

- Workloads
- Applications
- System size
- Network topology
- Energy consumption
- Scheduling policies

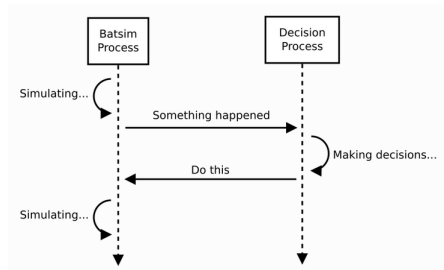
Complex implementations:
Kubernetes default scheduler
weighs **47k lines of code**.

Studying RJMS



Different approaches

- Analytical study
- Real experiments
- Emulation
- **Simulation**



Batsim, an infrastructure simulator aimed at studying RJMS.

Contribution



- Event based
- Own protocol
- Dilated time



kubernetes

- Constant API requests
- Own protocol
- (Real) machine time



Batkube



kubernetes

Batkube supports

- Any Go scheduler
- Any cluster size
- Resource requests
- Non parallel tasks

Literature review

Infrastructure simulators

	Grid	HPC	Cloud	P2P	Volunteer
SimGrid	✓				
GridSim	✓				
LogGOPSim		✓			
BigSim		✓			
CloudSim			✓		
GroudSim	✓		✓		
PeerSim				✓	
OverSim				✓	
SimBA					✓
SimBOINC					✓

Domain specific simulators



SimGrid

- Framework for building simulators
- Versatile, accurate and scalable
- 20 years of experience
- Simple analytical models

Simulators for the study of RJMS

Often *ad hoc* simulators

“Publish and perish” - Milian Poquet

Some active projects

- YARNSim
- SLURM simulator

	Scheduler	Platform	Job model
Accasim	Internal	Ad hoc	Static duration
Alea	Internal	GridSim	Static duration
Batsim	Custom Protocol	SimGrid	SimGrid models

Simulation of RJMS

Kubernetes cluster simulation

Kubernetes simulation projects

	joySim	k8s-cluster-simulator
Origin	private (JD.com)	student project
Availability	closed source	open source
Focus	service	batch processing
Scheduler	any	user implementation
Models	mock nodes	static job durations

Notable Kubernetes schedulers

- kube-batch



VOLCANO
Kubernetes Native Batch System




Kubeflow

- Poseidon (Firmament)
- kube-scheduler

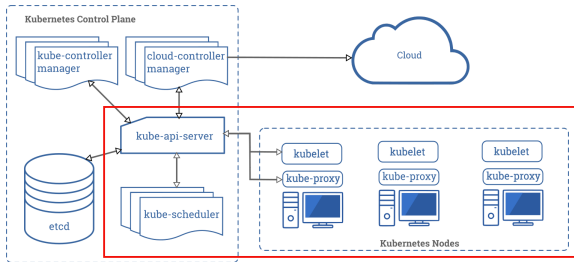
Integrating Kubernetes schedulers to Batsim

Different communication paradigms



source: <https://batsim.readthedocs.io>

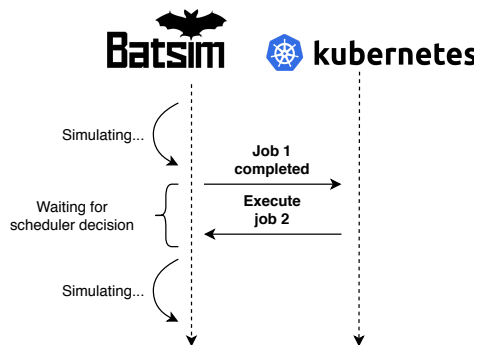
- Event based
- Simulation time




source: <https://kubernetes.io/docs/concepts/overview/components/>

- Central API
- Real time

Time synchronization



Scenario 1: correct synchronization



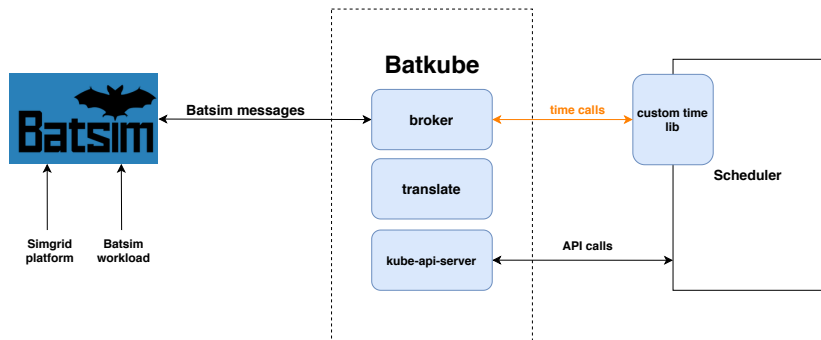
Scenario 2: delayed decision

Technical challenges

Challenges to tackle

- 1 Integration with Kubernetes
- 2 Scheduler time interception
- 3 Time synchronization

Architecture of Batkube

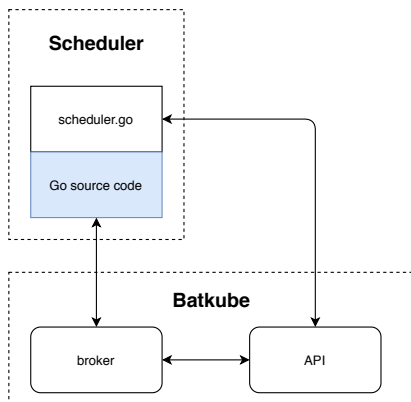


Global architecture of Batkube.

Time interception

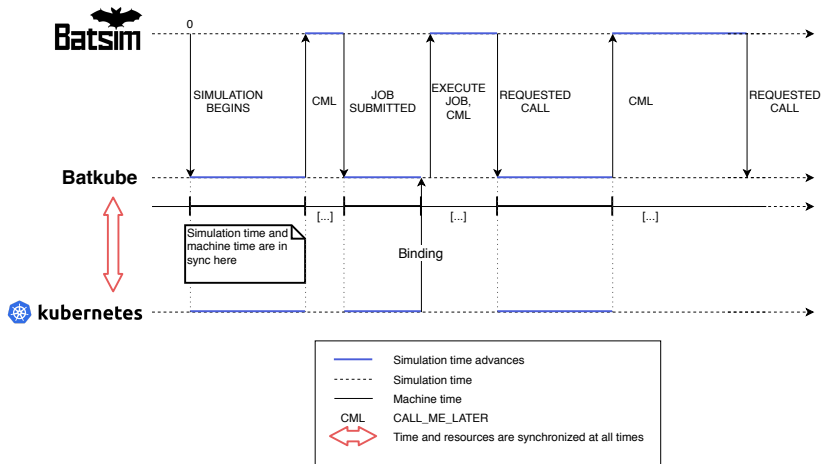
Redirection of time calls

- Specific functions are redirected
- Automatic source code manipulation using AST
- Ensured compatibility with the rest of the code



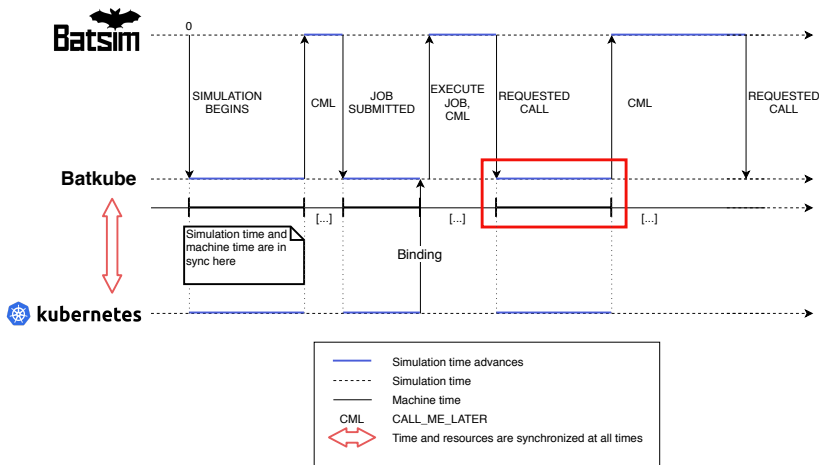
Schedulers are patched to redirect their time.

Time synchronization




Time synchronization between Batsim and the scheduler

Parameters of the synchronization I



Timeout value

Parameters of the synchronization II




Simulation time step $\in [\text{base-simulation-timestep}, \text{max-simulation-timestep}]$

Study of the simulator


Experimental design

TODO: Scheduler used, platforms and workloads tested, what experiments (parameters, metrics studied, repetitions)

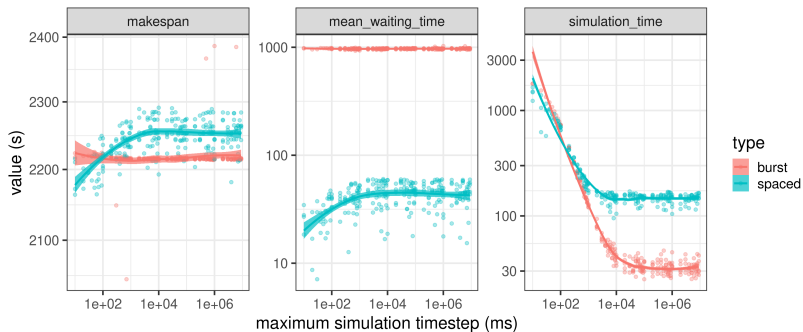
Timeout I




Timeout II



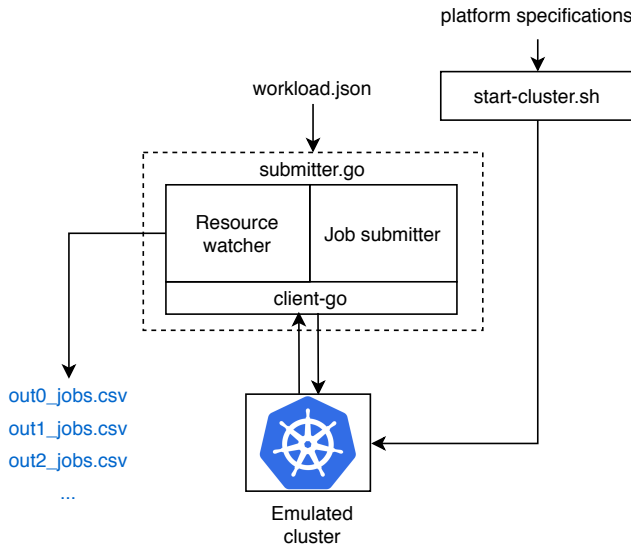
Maximum simulation timestep I



Maximum simulation timestep II



Experimentation on a real cluster



Deviation with reality

workload	makespan				mean waiting time			
	emulated		simulated		emulated		simulated	
	μ	σ	μ	σ	μ	σ	μ	σ
burst	2467	28.3	2215 (-252)	0.508	1077	10.6	970 (-107)	12.6
spaced	2468	5.14	2257 (-211)	16.9	146	1.67	48.1 (-97.9)	9.44
realistic	32556	-	32555 (-1)	1.30	2884	-	2020 (-864)	950

Conclusion

Deviation with reality: can be fixed with some work on the api. Need experiments to measure and quantify this deviation.

max timestep: studying max timestep alone is not enough, need to study it with backoff multiplier.

base time step: need an experiment on it. Too much importance was credited to max timestep, the base timestep might have importance.

Discussion and future work

Capabilities and limitations of Batkube

WIP

Capabilities

- Delay jobs
- Cpu and memory requests
- Can patch any kubernetes scheduler written in Go
- The api only supports the default scheduler

Limitations

- Memory hungry (in fact, the scheduler is memory hungry)
- Some problems with the scheduler
- Not scalable

Perspectives for future work


- parallel jobs
- storage
- more complete api: support for more schedulers but also tools (monitoring tools)

References I

Any questions?

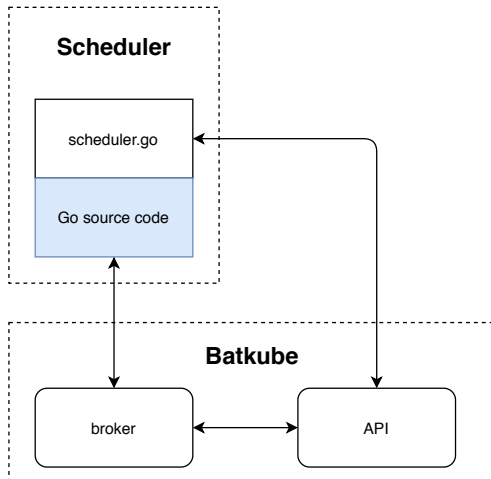
Thank you for your attention!
I am open to any questions.

Batkube integration with Kubernetes



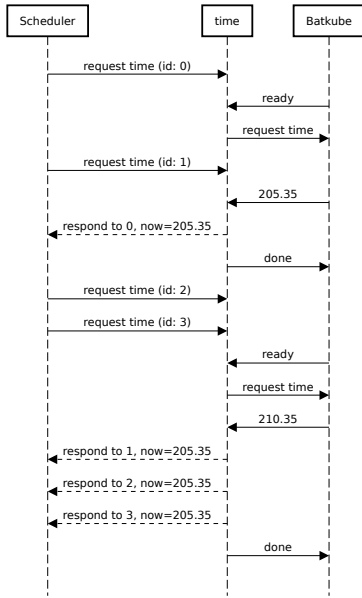
Reimplementation of a custom API.

Time interception




Schedulers are patched to redirect their time.

batsky-go



Exchanges between the scheduler, batsky-go (“time”) and Batsim

Similar resources



source: <https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>

Translation between Kubernetes and Batsim

- A Pod = a job.
- A Node = a compute resource.