

Development and evaluation of a Kubernetes cluster simulator based on Batsim

Presented by: Théo Larue
Supervised by: Olivier Richard & Michael Mercier

Université Grenoble Alpes

August 31, 2020



Table of contents

- 1 Introduction
- 2 Literature review
- 3 Integrating Kubernetes schedulers to Batsim
- 4 Study of the simulator
- 5 Discussion and future work


Introduction

Resource and Jobs Management System

The RJMS is at the core of the cluster.

Examples of RJMS

- OAR
- SLURM
- HadoopYARN
- Apache Mesos



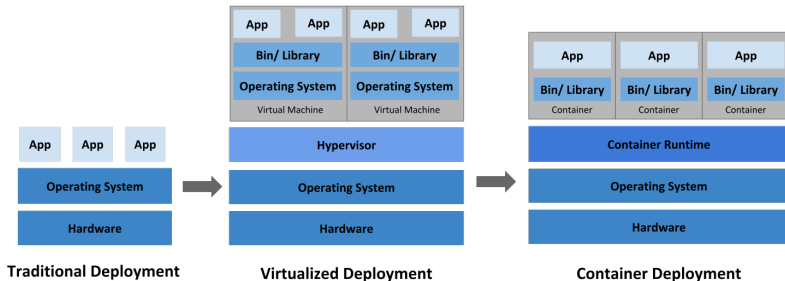
Kubernetes

Kubernetes in a nutshell


- Open source resource manager for containerized applications
- About 2M lines of code
- 2.8k contributors



kubernetes



A component of the RJMS: the scheduler



Scheduling is the act of allocating tasks to resources.

Numerous factors

- Workloads
- Applications
- System size
- Network topology
- Energy consumption
- Scheduling policies

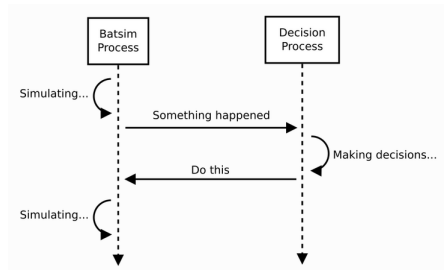
Complex implementations:
Kubernetes default scheduler
weighs **47k lines of code**.

Studying schedulers



Different approaches

- Analytical study
- Real experiments
- Emulation
- **Simulation**



Batsim, an infrastructure simulator aimed at studying RJMS.

Contribution



- Event based
- Own protocol
- Dilated time



kubernetes

- Constant API requests
- Own protocol
- (Real) machine time



Batkube



kubernetes

Batkube supports

- Any Go scheduler
- Any cluster size
- Resource requests
- Non parallel tasks

Literature review

Domain specific simulators

refs on domain specific simulators

Software specific simulators

YARNSim, SLURM simulator

Publication specific simulators

“Publish and perish” - Milian Poquet

SimGrid

SimGrid: Versatile, scalable, accurate.

Cpu = a computation speed.

Storage = a seek time and a data transfert rate.

Network = a flow model, modeling bandwidth sharing behaviors.

Simple models but thoroughly validated.

Related work

GridSim

Alea: modular, extensible.

Accasim: supports additional information (temperature, power consumption). Very efficient in terms of simulation time and memory usage.

Kubernetes cluster simulation

k8s-cluster-simulator: open source, student project, delay jobs.

Schedulers provided via a Go interface.

joySim: closed-source, fully fledged kubernetes cluster simulator, service oriented (mock nodes).

Integrating Kubernetes schedulers to Batsim

Technical challenges

Challenges to tackle

- 1 Integration with Kubernetes.

Technical challenges

Challenges to tackle


- 1 Integration with Kubernetes.
- 2 Intercepting scheduler time.

Technical challenges

Challenges to tackle

- 1 Integration with Kubernetes.
- 2 Intercepting scheduler time.
- 3 Time synchronization between Batsim and the scheduler.


Batsim concepts



source <https://batsim.readthedocs.io>

Batsim events and protocol.
User defined workloads.
(insert json examples?)


Kubernetes concepts



source: <https://kubernetes.io/docs/concepts/overview/components/>

Kubernetes components.

Kubernetes concepts



source: <https://kubernetes.io/docs/concepts/overview/components/>

Kubernetes components.

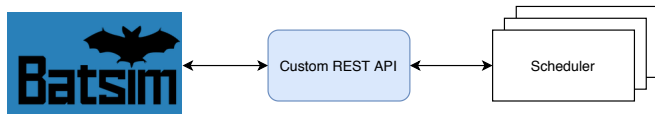
Different paradigms

Batsim: event based, simulation time.

Kubernetes scheduler: asynchronous calls to the API, machine time.

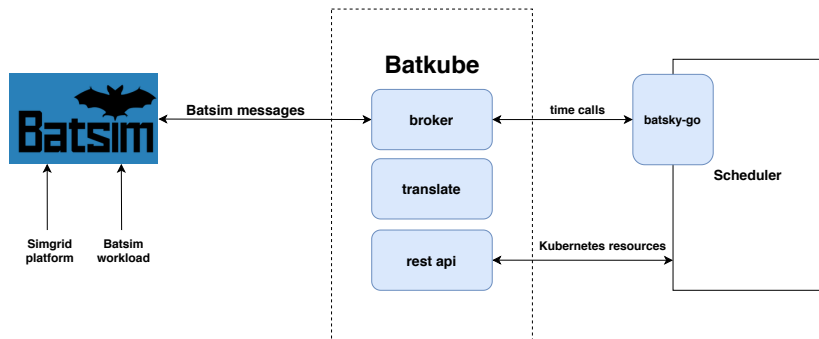
The goal is to make the scheduler event based and relying on simulation time for Batsim, and make Batsim a kube-api-server to the scheduler.

Batkube integration with Kubernetes



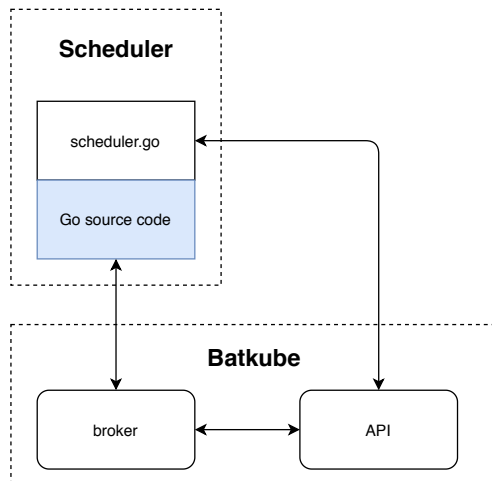
Reimplementation of a custom API.

Architecture of Batkube



Global architecture of Batkube.

Time interception

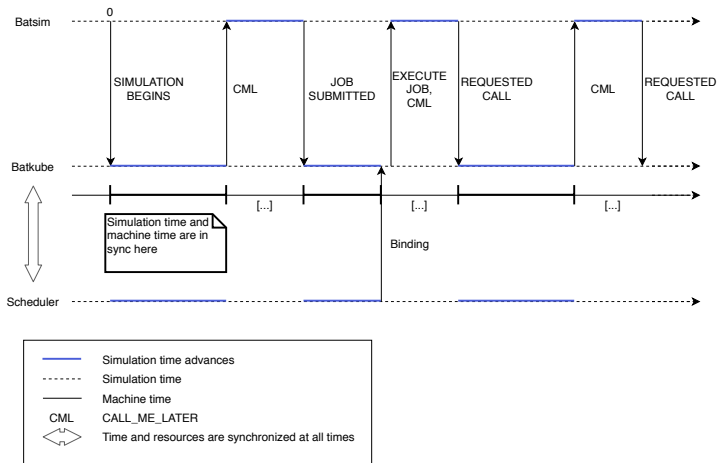


Schedulers are patched to redirect their time.

Time synchronization I


TODO: explain CML

Time synchronization II



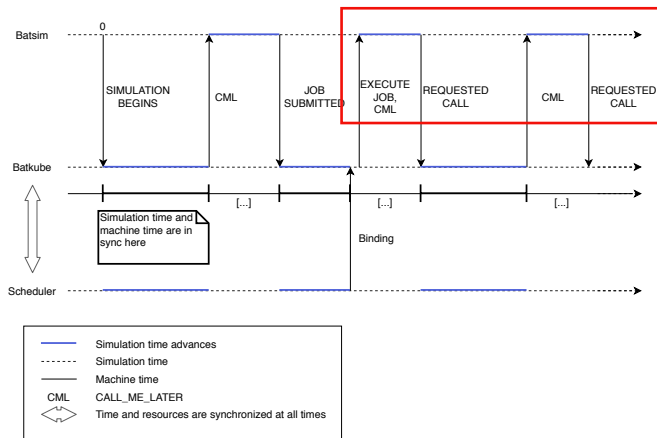
Time synchronization between Batsim and the scheduler

Parameters of the synchronization I




Timeout value

Parameters of the synchronization II



Simulation time step $\in [\text{base-simulation-timestep}, \text{max-simulation-timestep}]$

Time synchronization breakdown



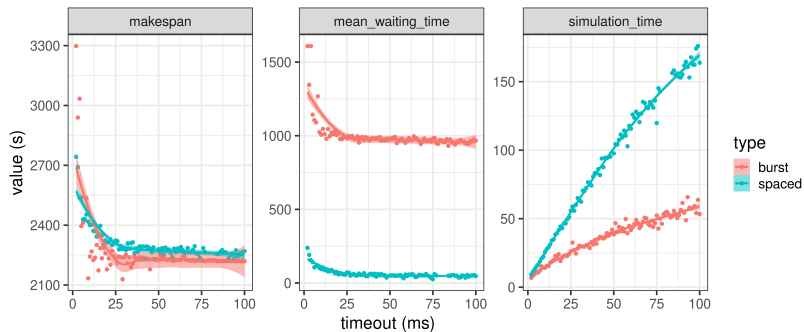
Time synchronization between Batsim and the scheduler

Study of the simulator


Experimental design

TODO: Scheduler used, platforms and workloads tested, what experiments (parameters, metrics studied, repetitions)


Timeout I




Timeout II



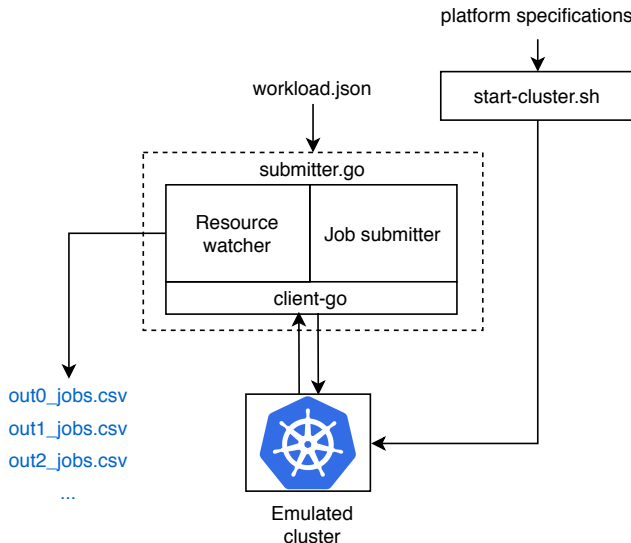
Maximum simulation timestep I



Maximum simulation timestep II



Experimentation on a real cluster



Deviation with reality

workload	makespan				mean waiting time			
	emulated		simulated		emulated		simulated	
	μ	σ	μ	σ	μ	σ	μ	σ
burst	2467	28.3	2215 (-252)	0.508	1077	10.6	970 (-107)	12.6
spaced	2468	5.14	2257 (-211)	16.9	146	1.67	48.1 (-97.9)	9.44
realistic	32556	-	32555 (-1)	1.30	2884	-	2020 (-864)	950

Conclusion

Deviation with reality: can be fixed with some work on the api. Need experiments to measure and quantify this deviation.

max timestep: studying max timestep alone is not enough, need to study it with backoff multiplier.

base time step: need an experiment on it. Too much importance was credited to max timestep, the base timestep might have importance.

Discussion and future work

Capabilities and limitations of Batkube

WIP

Capabilities

- Delay jobs
- Cpu and memory requests
- Can patch any kubernetes scheduler written in Go
- The api only supports the default scheduler

Limitations

- Memory hungry (in fact, the scheduler is memory hungry)
- Some problems with the scheduler
- Not scalable

Perspectives for future work

- parallel jobs
- storage
- more complete api: support for more schedulers but also tools (monitoring tools)

References I